

Level and Shape: Learning Adjacent-Lanthanide Selectivity from the Solvent-Extraction Record

Bogdan Mironov

Dias Daulet

Department of Chemistry, Berea College
Berea, KY, USA

NIS Nauryzbay
Almaty, Kazakhstan

`mironovb@berea.edu`

`daulet_d0806@hbalm.nis.edu.kz`

Konstantinos D. Vogiatzis

Department of Chemistry, University of Tennessee
Knoxville, TN, USA

`kvogiatz@utk.edu`

September 12, 2026

Abstract

Adjacent lanthanide ions differ in radius by about 0.013 Å, and separating them is the most expensive step in rare-earth processing. In this work, we show that the published solvent-extraction record contains sufficient information to learn which ligands perform this separation effectively. On 4,746 measurements covering 162 extractants, a gradient-boosting model evaluated with whole extractants held out predicts log separation factors that correlate with experiment at $r = 0.59$ across 905 adjacent pairs. The prediction is split in two: one model estimates the overall extraction level of each series, the other the differences between its metals, and the pair outperforms a single model trained directly on the separation factors. Used as a screen, the model identifies the most effective separations (above a factor of two) with an AUC of 0.81, and extractants ranked in the model’s top quartile fall within the true top quartile 50 % of the time, twice the rate expected by chance. The conformal prediction intervals match their nominal coverage. A graph neural network operating on three-dimensional metal–ligand geometries provides similar predictions independently complex and adds a small further gain in combined model.

1 Introduction

Trivalent lanthanides scarcely differ chemically. Along the series, the ionic radius decreases by about 0.013 Å from one lanthanide to the next [1]. A ligand that separates neodymium from praseodymium has, in effect, mostly only that shrinkage to use for the separation. An industrial plant therefore runs its extraction stages in long series. New extractants still found, for the most part, from trial and error [2].

The SAFE database [3] compiles distribution ratios from the solvent-extraction literature, and many of its ligands have been measured across most of the series under a single set of conditions. The separation factor of a pair of metals is the ratio of their distribution ratios from the same experiment, so errors that are common to both measurements cancel. This paper addresses how much adjacent-lanthanide selectivity can be learned from this record for extractants not seen during training, and what a three-dimensional (3D) structure of the metal–ligand complex adds. The work makes the following three contributions: (i) an evaluation score and a model that both separate extraction strength from selectivity ($+0.27 \rightarrow +0.33$ from the model split with the same learner and features); (ii) a small gain from 3D structures of the complexes; and (iii) a benchmark

of three tight-binding Hamiltonians on the lanthanide contraction, which measures how much information about the metal a computed geometry contains. Past a few dozen training extractants, the learning curve levels off.

2 Data and target

From the SAFE Database lanthanide(III) rows with a resolvable extractant structure were kept. We canonicalise the SMILES, remove duplicates and detection-limit placeholders, and retain the 4,746 measurements of $\log D$ on 162 extractants for which a complex geometry exists. Each row carries the extractant, the metal and 64 condition columns. The 956 complexes were built with Architector [4] (extractant and nitrate or water to coordination number 8–10) and optimised with GFN2-xTB [5].

A separation factor is only defined between measurements made under the same conditions, therefore rows are grouped into blocks (one extractant, one bin of each of ten condition descriptors). Within a block, replicate $\log D$ measurements of one metal are averaged into a cell, and each pair of cells whose lanthanides lie adjacent in the series contributes one target value,

$$\log \text{SF}_{i,i+1} = \log D_{\text{Ln}_i} - \log D_{\text{Ln}_{i+1}}. \quad (1)$$

Models are trained on $\log D$ row by row and evaluated after this construction is applied to their out-of-fold predictions; the score R_{adj}^2 is the coefficient of determination of predicted against measured $\log \text{SF}$ over all adjacent pairs.

Every number presented below is out-of-fold under five-fold cross-validation grouped by extractant, repeated three times. The population yields 905 adjacent pairs on 76 extractants.

3 Model

The tabular representation follows Liu et al. [2]: RDKit descriptors and 2048-bit Morgan fingerprints of the extractant [6], the metal’s atomic number, position in the series and Shannon radius, and the 64 condition columns, 746 in total. The learner is CatBoost [7] with a quantile loss ($\alpha = 0.6$), depth 9 and feature subsampling 0.3, chosen on the adjacent-pair score. Trained on $\log D$ directly it reaches $R_{\text{adj}}^2 = +0.268$.

The raw target combines the overall extraction strength of a ligand with its variation across metals. Only the variation is scored, so each part receives its own model. The level model f_L is trained on $\log D$, and only its test-block mean, the anchor a_b , is used. The shape model f_S is trained on the block-centred target $\log D_i - \bar{y}_{b(i)}$, and its predictions are centred within the test block in the same way,

$$\hat{y}_i = a_{b(i)} + \left(f_S(x_i) - \text{mean}_{j \in b(i)} f_S(x_j) \right). \quad (2)$$

For the 3D channel, heavy atoms of the complex are nodes, atoms within 4 Å are connected, and three continuous-filter layers of width 96 [8] feed a readout with a metal-centred term and a radial shell histogram, since plain mean pooling averages away shifts of a few hundredths of an Ångström; the embedding is concatenated with the tabular features, and training uses a Huber loss on $\log D$ plus a contrast term on the within-block pairs of each batch (Appendix A). On its own the network scores +0.266, no better than the boosting baseline. Its block-centred prediction $s_{3\text{D}}$ enters Eq. (2) as a second shape source, $\hat{y} = a + (1 - w) s_{\text{tab}} + w s_{3\text{D}}$, with $w = 0.35$ chosen on the training folds only; the score becomes +0.326.

4 Results

The correlation is robust and ligand-specific. On held-out extractants, predicted and measured separation factors correlate with $r = 0.59$ (Figure 1a). A cluster bootstrap over extractants gives an interval of +0.21 to +0.38 for R_{adj}^2 . As a null test, the predictions were shuffled among extractants within each pair position. The shuffled predictions score $R_{\text{adj}}^2 = -0.05$. Removing both the position means and the extractant means (10 and 13 % of the variance) leaves the ligand-by-position interaction, which is the selectivity itself. On this term the model still scores +0.23.




Figure 1: (a) Predicted against measured adjacent-pair logSF for the 905 pairs, extractants held out; the dashed line is the identity. (b) Per adjacent pair: Spearman ρ between predicted and measured extractant rankings, and the top-quartile hit rate against chance. (c) Slope of the Ln-donor distance against the Shannon radius (the compliance c_L) for the 71 ligands under three Hamiltonians: points are ligands, shaded shapes their kernel density, bars the mean; 1 is exact agreement with the radii. (d) Learning curve: small points are the 52 individual runs (four CV seeds, three random draws of the training extractants) at their own training size, large markers the mean per fraction; solid and dotted lines the hyperbolic and power-law extrapolations; dashed line the full training set.

Screening performance. Ten percent of the pairs account for 71 % of the variance, so R_{adj}^2 mostly reflects the large separations. The questions that matter in screening are therefore evaluated directly (Figure 1b). For a given adjacent pair, ranking the held-out extractants gives a mean Spearman correlation of 0.46 over the twelve positions, and the predicted top quartile contains a measured top-quartile ligand 50 % of the time, against 25 % by chance. Used to detect separations with $\text{SF} \geq 2$ (base rate 0.20), the same predictions give an AUC of 0.81. Split-conformal intervals [9], calibrated leave-one-extractant-out, cover 79 % and 90 % of the measured values at nominal levels of 80 and 90 %.

Table 1: Adjacent-pair log SF R^2 on the 905 pairs, extractants held out. Pair identity alone scores +0.066.

Model	R^2
CatBoost on log D	+0.268
3D encoder, standalone	+0.266
Level + shape, Eq. (2)	+0.318
Level + shape + 3D shape, $w = 0.35$	+0.326

5 Contribution of three-dimensional structure

The optimised complexes enter the model through the shape term only. Adding the block-centred prediction of the 3D encoder to the tabular shape with weight $w = 0.35$ raises the score from +0.318 to +0.326 (eight CatBoost seeds; 32-seed encoder ensemble), whereas the same geometries converted to 84 descriptor columns and given to a model that predicts log D directly lower it from +0.278 to -0.022 . The gain lies in the light half of the series: on the 482 La–Gd pairs the score rises from +0.356 to +0.369, while on the 347 Tb–Lu pairs it remains at +0.028. The heavy half is where every model has little skill (Spearman 0.30 against 0.58) and where the separations are smaller (sd 0.17 against 0.32).

The gain is limited by the geometries, which vary with the metal through a single number. Eight encoders with different architectures, cutoffs and losses give nearly identical predictions ($r = 0.990$ between disjoint eight-seed halves; effective rank 1.05 of the eight prediction vectors). In GFN2-xTB the lanthanide parameters are linear interpolations in Z between Ce and Lu, so the optimised metal–donor distances follow the ionic radius, which the tabular features already contain. The contraction benchmark (Figure 1c; 71 ligands, fifteen lanthanides substituted into one relaxed complex per ligand and re-optimised with each Hamiltonian) gives a slope of the mean metal–donor distance against the Shannon radius of 0.405 ± 0.145 for GFN2-xTB, 1.078 ± 0.094 for g-xTB [10], and 1.081 ± 0.214 for Ln-xTB [11], whose series is irregular (linear-fit R^2 0.43 against 0.82 for g-xTB). Replacing the GFN2-xTB geometries with g-xTB geometries changes the score by +0.005 ($p = 0.68$, eight paired seeds; Appendix B), and the ligand-specific part of the slope is uncorrelated with measured selectivity under all three Hamiltonians ($r = 0.11, -0.02, 0.17$; 44 ligands).

How much data the problem needs. Seventy-six extractants contribute scored pairs, so the obvious worry is that the model is starved. It is not: trained on a random fraction of each fold’s extractants, test folds untouched, the score climbs from +0.15 with eight extractants to +0.30 with forty-six and then flattens at +0.31 with all sixty-one (Figure 1d). The last third of the data is worth +0.005, inside seed noise; the two fitted extrapolations put the gain from doubling the extractant count between +0.01 and +0.04. A few dozen well-measured ligands are enough; what is missing is what separates ligands within one class (Appendix D).

Limitations and conclusion. The 3D gain is small against seed variability, and the structure set omits, by construction, the quantities that could differ between ligands: the coordination number follows a rule of the metal alone, 89 % of the complexes are the bare Ln^{3+} cation rather than the extracted species, and only geometries, not binding energies, are used. These are the natural next inputs for the shape term, and the compliance benchmark is a cheap check of whether a candidate structure set carries more than the ionic radius. No ceiling is identified: separations reproduce to about 0.16 on a spread of 0.22, but 70 % of scored pairs have no replicate, and within one extractant the model’s series shape is weak (median per-extractant $R^2 = 0.09$). What it predicts well is which ligands are selective and where their large separations lie. The step that mattered was asking the model the right question: the shape of a series, not its level.

References

- [1] R. D. Shannon. Revised effective ionic radii and systematic studies of interatomic distances in halides and chalcogenides. *Acta Crystallographica Section A*, 32(5):751–767, 1976. doi: 10.1107/S0567739476001551.

- [2] Tongyu Liu, Katherine R. Johnson, Santa Jansone-Popova, and De-en Jiang. Advancing rare-earth separation by machine learning. *JACS Au*, 2(6):1428–1434, 2022. doi: 10.1021/jacsau.2c00122.
- [3] L. da Silva Garcia Leite, B. Zhang, Z. Acar, J. Elowitz, L. J. Augustine, A. E. Clark, V. Karamalis, D. Perez, J. Schrier, M. Taylor, and P. Yang. Creation of the separation archive for elements (SAFE) database. *Solvent Extraction and Ion Exchange*, 43(7):847–852, 2025. doi: 10.1080/07366299.2025.2564381.
- [4] Michael G. Taylor, Daniel J. Burrill, Jan Janssen, Enrique R. Batista, Danny Perez, and Ping Yang. Architector for high-throughput cross-periodic table 3D complex building. *Nature Communications*, 14: 2786, 2023. doi: 10.1038/s41467-023-38169-2.
- [5] Christoph Bannwarth, Sebastian Ehlert, and Stefan Grimme. GFN2-xTB — an accurate and broadly parametrized self-consistent tight-binding quantum chemical method with multipole electrostatics and density-dependent dispersion contributions. *Journal of Chemical Theory and Computation*, 15(3):1652–1671, 2019. doi: 10.1021/acs.jctc.8b01176.
- [6] David Rogers and Mathew Hahn. Extended-connectivity fingerprints. *Journal of Chemical Information and Modeling*, 50(5):742–754, 2010. doi: 10.1021/ci100050t.
- [7] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. CatBoost: Unbiased boosting with categorical features. In *Advances in Neural Information Processing Systems*, volume 31, pages 6638–6648. Curran Associates, Inc., 2018. URL <https://proceedings.neurips.cc/paper/2018/hash/14491b756b3a51daac41c24863285549-Abstract.html>.
- [8] Kristof T. Schütt, Pieter-Jan Kindermans, Huziel E. Sauceda, Stefan Chmiela, Alexandre Tkatchenko, and Klaus-Robert Müller. SchNet: A continuous-filter convolutional neural network for modeling quantum interactions. In *Advances in Neural Information Processing Systems*, volume 30, pages 991–1001. Curran Associates, Inc., 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/303ed4c69846ab36c2904d3ba8573050-Abstract.html>.
- [9] Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005. doi: 10.1007/b106715.
- [10] Thomas Froitzheim, Marcel Müller, Andreas Hansen, and Stefan Grimme. g-xTB: A general-purpose extended tight-binding electronic structure method for the elements H to Lr ($Z = 1\text{--}103$). ChemRxiv, 2025. URL <https://doi.org/10.26434/chemrxiv-2025-bjxvt>.
- [11] Yang-Yang Zhang. New optimized parameters of extended tight-binding method Ln-xTB (Ln = La–Lu) to explore lanthanide molecular chemistry. *Journal of Computational Chemistry*, 47(4):e70321, 2026. doi: 10.1002/jcc.70321.

A Protocol

Code, out-of-fold predictions and every negative result: <https://github.com/mironovb/lanthanidestrain>.

Blocks. The ten binned condition descriptors that define a block are acid class, acid strength, nitrate activity, pH, diluent family, modifier class, temperature, contact time, phase ratio and metal concentration. Pairs with a measured logSF of exactly zero, which arise when both metals sit at a detection limit, are excluded. Promethium is absent, so Nd–Sm is not an adjacent pair.

Hyperparameters. CatBoost: 1,500 iterations, learning rate 0.04, depth 9, ℓ_2 regularisation 3, feature subsampling 0.3, quantile loss $\alpha = 0.6$, identical for the level and shape models. 3D encoder: width 96, three continuous-filter layers, 64 radial basis functions to 5 Å, cutoff 4 Å, heavy atoms only, dropout 0.15, head width 256, Adam at 2×10^{-3} with weight decay 10^{-4} , 60 epochs with early stopping on a validation fold, batches of 64 rows grouped by block so the contrast term always has pairs, pair-contrast weight 4 with adjacent pairs weighted $3\times$. Blend weight: for each fold, selected over $\{0, 0.05, \dots, 1\}$ on the other folds’ out-of-fold pairs, with equal weight per extractant.

Learning curve. Fractions 0.125, 0.25, 0.5, 0.75 and 1 of each fold’s training extractants; within a draw, the subset for a smaller fraction is contained in the subset for the larger one. The full-data run reproduces the stored champion exactly. Because the anchor is one number per extractant it cancels in every adjacent pair, so the level model does not affect the score and only the shape model’s training set matters.

B Controls

Interventions under the same protocol on the 905 pairs; ΔR^2 relative to the best comparable system.

Intervention	ΔR^2
Simplicial encoder instead of the distance graph	0.00 ($r = 0.96$, blend weight 0.01)
Persistence descriptors as extra shape features	−0.03 to −0.09
Metal-feature weight $\times 2$ / $\times 5$ / $\times 10$	+0.00 / −0.05 / −0.08
All-pairs difference model added to the blend	0.00
Condition-transfer features	+0.004
g-xTB geometries instead of GFN2-xTB (802 complexes)	+0.005 ($p = 0.68$)
Complexes rebuilt in correspondence	−0.013
Solvent-relaxed geometries, water / octanol	−0.008 / −0.081
GFN2 / g-xTB interaction-energy features	0.00

C Ln-xTB parameters

The Ln-xTB parameters are given in the supporting information of Zhang [11] as PDF tables only. We parsed both copies (Tables S60–S61 and S62; no disagreement), replaced the $Z = 57$ –71 blocks of the stock GFN2-xTB parameter file, and ran unmodified xtb 6.7.1. Single points on the 35 optimised structures shipped with the paper reproduce the published energies to better than $2 \times 10^{-8} E_h$ once the paper’s spin convention is used (the high-spin count of unpaired f electrons, reduced by one where its parity is inconsistent with the electron count). All 2,130 optimisations of the benchmark converged. A closed-shell control gives $c_L = 0.981 \pm 0.140$ with the same jagged series.

D Americium

SAFE holds 1,794 Am(III) rows on 110 ligands, 95% of them ligands with lanthanide data, giving 312 Am/Eu separation factors on 82 extractants under identical conditions. Zero-shot, with Am treated as a lanthanide at its radius, the Am/Eu separation is not predicted ($R^2 = -0.14$). Joint training with the Am rows reaches +0.47 on held-out extractants, but a control that assigns each held-out extractant the mean of the other extractants in its ligand family scores +0.47 as well, and within a family the model has no skill (+0.02); the lanthanide score is unchanged ($+0.002 \pm 0.020$ over paired seeds). The model learns which ligand classes discriminate Am from Eu, nothing finer.